

BOĞAZIÇI UNIVERSITY

DEEP LEARNING
CMPE 597

Assignment 1

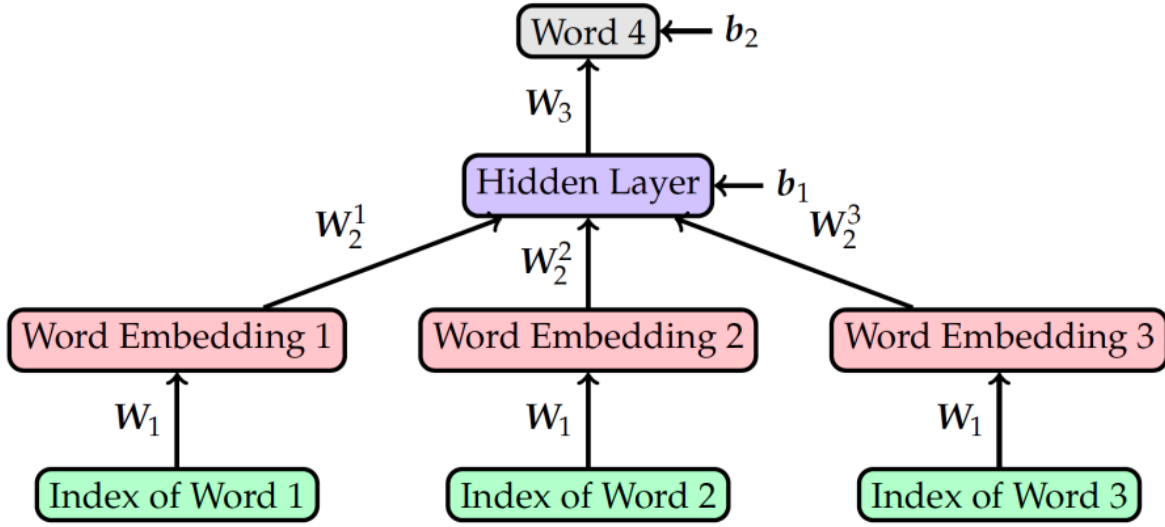
Author:

Y. Harun Kıvrıl
02 May 2021



Introduction

This project aims to implement the following feed forward neural network structure by only using the libraries for calculations. After the implementation the network is trained on train data with hyperparameter tuning on validation data. The network obtains the best validation accuracy is selected and it evaluated on test data. For the evaluation besides the accuracy and loss values a tsne plot of the embeddings are investigated and the model's predictions for three examples are evaluated.



1 Forward Propagation

Forward propagation equation of the given network with mini batch size n and dictionary size 250 can be written in the following form.

Let $I_i \in \mathbb{R}^{n \times 250}$ be the one hot encoded batch of i th input word and W_1 be the weight matrix of embedding layer. Then the embedding of the input words are:

$$e^{[1]} = I_1 W_1, \quad e^{[2]} = I_2 W_1, \quad e^{[3]} = I_3 W_1$$

where $W_1 \in \mathbb{R}^{250 \times 16}$ and as a result $e^{[i]} \in \mathbb{R}^{n \times 16}$. After that each embedding are concatenated into a single layer:

$$e = [e^{[1]}, e^{[2]}, e^{[3]}]$$

where $e \in R^{nx48}$. Later this embeddings are mapped to the hidden layer using sigmoid activation function σ_1 :

$$h = eW_2 + b_1$$

$$o^{[1]} = \sigma_1(h)$$

where $W_2 \in R^{48 \times 128}$, $b_1 \in R^{128}$, h and $o^{[1]} \in R^{nx128}$. Finally the outputs of the hidden layer is mapped to the output layer using softmax activation function σ_2 :

$$z = o^{[1]}W_3 + b_2$$

$$o^{[2]} = \sigma_2(z)$$

where $W_3 \in R^{128 \times 250}$, $b_2 \in R^{250}$, and $o^{[2]}, z \in R^{nx250}$.

After obtaining the outputs from network, the cross entropy loss is calculated using the following formula:

$$L = -\frac{1}{n} \sum_{i=1}^n t_i \log(o_i^{[2]})$$

where $t_i \in R^{250}$ is the true label of i th observation in mini batch and $o_i^{[2]}$ is the model output of that observation.

2 Backward Propagation

The back propagation of the given network can be written in the following form using the chain rule. In order to do this an intermediate variable δ is introduced to store common terms. For the derivative of W_3 and b_2 :

$$\delta_3 = \frac{\partial L}{\partial o^{[2]}} \frac{\partial o^{[2]}}{\partial z}$$

$$\frac{\partial L}{\partial W_3} = \delta_3 \frac{\partial z}{\partial W_3}$$

$$\frac{\partial L}{\partial b} = \delta_3 \frac{\partial z}{\partial b_2}$$

When we calculate this chain one by one:

$$\frac{\partial L}{\partial o_i^{[2]}} = \frac{-t_i}{o_i^{[2]}}$$

and

$$\frac{\partial o_i^{[2]}}{\partial z_j} = \begin{cases} o_i^{[2]}(1 - o_j^{[2]}) & i = j \\ -o_i^{[2]}o_j^{[2]} & i \neq j \end{cases}$$

where o_i is the result of the i^{th} output neuron and z_j is the output of the j^{th} neuron before activation. When the derivatives of cross entropy loss and softmax activation are combined the result becomes $o^{[2]} - t$ and this result can be used to make calculations faster. The network has both implementations but uses the second one to obtain results faster. Additionally:

$$\frac{\partial z_i}{\partial W_{3_j}} = o^{[1]} \quad \frac{\partial z}{\partial b_2} = 1$$

For the derivative of W_2 and b_1 :

$$\delta_2 = \delta_3 \frac{\partial z}{\partial o^{[1]}} \frac{\partial o^{[1]}}{\partial h}$$

$$\frac{\partial L}{\partial W_2} = \delta_2 \frac{\partial h}{\partial W_2}$$

and for bias:

$$\frac{\partial L}{\partial b_1} = \delta_2 \frac{\partial h}{\partial b_1}$$

In order to calculate these gradients we need to the following calculations:

$$\frac{\partial z^{[2]}}{\partial o^{[1]}} = W_3$$

$$\frac{\partial o^{[1]}}{\partial h} = o^{[1]}(1 - o^{[1]})$$

and

$$\frac{\partial h_j}{\partial W_{2_i}} = e \quad \frac{\partial z}{\partial b^1} = 1$$

For derivative of W_1 :

$$\delta_1 = \delta_i \frac{\partial h}{\partial e}$$

Since, there are three words as input and they all use W_1 , δ_1 is needed to be separated according to corresponding $e^{[i]}$ parts:

$$[\delta_{1,1}, \delta_{1,2}, \delta_{1,3}] = \delta_1$$

Then the derivative becomes:

$$\frac{\partial L}{\partial W_1} = \sum_{i=1}^3 \delta_{1,i} \frac{\partial e^{[i]}}{\partial W_1}$$

$$\frac{\partial h}{\partial e} = W_2$$

$$\frac{\partial e^{[i]}}{\partial W_1} = I_i$$

Since, the training done in minibatches for gradient estimations, the total gradient obtained from the minibatch is divided to the batch size. As a summary the gradients are calculated in the following way:

$$\delta_3 = o^{[2]} - t \in R^{n \times 250}$$

$$\frac{\partial L}{\partial W_3} = \frac{1}{n} [\delta_3^T o^{[1]}]^T \in R^{128 \times 250}$$

$$\frac{\partial L}{\partial b_2} = \frac{1}{n} \sum_{i=0}^n \delta_{3_i} \in R^{250}$$

$$\delta_2 = (\delta_3 W_3^T) \odot (o^{[1]}(1 - o^{[1]})) \in R^{n \times 128}$$

$$\frac{\partial L}{\partial W_2} = \frac{1}{n} [\delta_2^T e]^T \in R^{48 \times 128}$$

$$\frac{\partial L}{\partial b_1} = \frac{1}{n} \sum_{i=0}^n \delta_{2_i} \in R^{128}$$

$$\delta_1 = (\delta_2 W_2^T) \in R^{n \times 48}$$

$$[\delta_{1,1}, \delta_{1,2}, \delta_{1,3}] = \delta_1$$

$$\frac{\partial L}{\partial W_1} = \sum_{j=0}^3 \frac{1}{n} \sum_{i=0}^n [\delta_{1,i,j}^T I_{j_i}]^T \in R^{250 \times 16}$$

3 Training and Evaluating the Network

The architecture given in the figure is implemented in a network class using python's numpy[2] library. The weight and bias matrices are initialized randomly using normal distribution with $\mu = 0$ and $\sigma = 0.01$. The forward and backward functions of the network are written according to the equations given above. Also a function for updating the weights with given learning rate and network's gradients is introduced.

Since train, validation and test sets are already separated, the only thing missing for training was the hyperparameters: learning rate and batch size. In order to tune this parameters three candidate levels are selected as large, medium and small. For the learning rate (1, 0.1, 0.01) are selected and for batch size (16, 128, 512) is considered. Then the network is trained with each combination of these values and the train and validation loss and accuracies are recorded for each epoch. Since it

takes a considerable amount of time to train the network, the maximum epoch number is determined as 250 and an early stopping mechanism is introduced. This mechanism stopped the training if the network obtained less validation accuracy than the best accuracy so far in 8 consecutive epochs. The selection of the hyperparameters is made based on the epoch with best validation accuracy in each hyperparameter combination. At the end of the tuning procedure the following table is obtained:

LR/Batch Size	Validation Accuracy		
	16	128	512
0.01	0.3683	0.2129	0.2126
0.1	0.364	0.366	0.3577
1	0.3384	0.3575	0.3698

Figure 1: Results obtained from hyperparameter tuning

According to these results. The batch size is selected as 512 and the learning rate is chosen as 1. There are some accuracies far away from the maximum one. The reason for that is the updates become smaller with decreasing learning rate and increasing batch size (since we use the average of the batch as estimate) and 250 epoch is not enough for them to give better results.

The following plots show the evolution of losses and accuracies in the training of the model with the selected hyperparameters. The plots are generated using python's matplotlib[3] library.




Figure 2: Training and validation loss per epoch for the selected model




Figure 3: Training and validation accuracy per epoch for the selected model

The selected model is also evaluated in the test set and following evaluation table is constructed.

Selected Model	Loss	Accuracy
Train	2.492	0.3818
Validation	2.6243	0.3698
Test	2.6268	0.3674

Figure 4: Training and validation loss per epoch for the selected model

4 Results and Discussion

As the table above shows, the model obtained 0.3674 accuracy on the test set which is similar to the validation performance. It is not a good ratio however while evaluating this ratio the following facts should be considered: the number of classes are very much and the model is a very simple language model and its architecture is given in advance.

Besides to accuracy evaluating the embeddings obtained from the model gives us an idea about the performance of the model. If the model has learnt something during the training we expect to have the similar words are close to each other in the embedding space. Since the embedding space is 16 dimensional, it is not possible to show the distances visually, therefore a dimensionality reduction technique called TSNE is used to reduce the dimension to 2.

The embeddings are obtained by multiplying the identity matrix of size 250 with the weights of the embedding layer W_1 . Then the embeddings are fed into sklearn's[5] TSNE[6] model and 2D representation of these embeddings are obtained. Later the embeddings are plotted with using the words as annotation.




Figure 5: TSN-e plot of the embeddings.

At the right bottom, there exists a cluster of auxiliary verbs such as might, could etc. Also, at the bottom middle it is possible to see the cluster of different forms of to be verb. At the right, there is a cluster of verbs and around (-5,-5) a cluster of numerical quantities can be found. Another cluster is around (-8,4) which includes the punctuation marks. Just on the right of it about (0,5) there is a cluster of words about the time such as ago, year, day etc. When the middle of the plot is investigated, there seems like a mess but it is possible to catch close words with close meanings. As a result it is possible to say that the embeddings obtained from the model makes sense.

When the prediction of the model on "city of new", "life in the", "he is the" is checked following results are obtained.

"city of new york" with prob: 0.9917

which absolutely makes sense.

"life in the game" with prob:0.1148

which also makes sense however not obvious as previous example. But it appears that there is a movie called "Life in the game" and the model can be considered successful in this task.

"he is the best" with prob: 0.2006

which could be a answer from a human being.

Therefore it is possible to conclude that the model is doing a good job on all of these examples. Following figure also shows the top 5 predictions and their probabilities to these sentences.

```
city of new :
Words , Probs
[['york' '0.9917425878583002']
 ['work' '0.002727040984006493']
 ['music' '0.0006320016791027593']
 ['.' '0.0005082617374240455']
 ['?' '0.0004943476940786268']]

life in the :
Words , Probs
[['game' '0.11480673123825035']
 ['country' '0.11018087247029701']
 ['world' '0.08720927304863811']
 ['united' '0.08375972719120026']
 ['city' '0.08239239385489469']]

he is the :
Words , Probs
[['best' '0.2006268697433097']
 ['first' '0.09665412542690183']
 ['only' '0.08004241887331866']
 ['same' '0.059552475088556385']
 ['law' '0.04427699275044061']]
```

Figure 6: Top 5 predictions and probabilities to examples

This table also shows that when there is not much options that make sense model is confident

and gives high probability to the best prediction. When it gives low probability it appears that there are also other options that may make sense.

5 Conclusion

In this project a deep language model with given architecture that takes three words as input and predicts the fourth word is implemented. After the implementation batch size and learning rate is tuned and the model with selected hyperparameters is evaluated. The model obtained 0.3674 accuracy on the test set similar to its validation accuracy. Even though the accuracy is low, when the embeddings are plotted in 2D, it is possible to see that meaning clusters occur. Also for the example inputs the model gives sensible results. Given these information model is considered to be a good model for the given architecture.

References

- [1] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [2] Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, Julian Taylor, Sebastian Berg, Nathaniel J. Smith, Robert Kern, Matti Picus, Stephan Hoyer, Marten H. van Kerkwijk, Matthew Brett, Allan Haldane, Jaime Fernández del Río, Mark Wiebe, Pearu Peterson, Pierre Gérard-Marchant, Kevin Sheppard, Tyler Reddy, Warren Weckesser, Hameer Abbasi, Christoph Gohlke, and Travis E. Oliphant. Array programming with NumPy. *Nature*, 585(7825):357–362, September 2020.
- [3] J. D. Hunter. Matplotlib: A 2d graphics environment. *Computing in Science & Engineering*, 9(3):90–95, 2007.
- [4] Michael A. Nielsen. Neural networks and deep learning, 2018.
- [5] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [6] Laurens Van der Maaten and Geoffrey Hinton. Visualizing data using t-sne. *Journal of machine learning research*, 9(11), 2008.

6 Appendix

6.1 Appendix A (Codes)

network.py

```
1 import numpy as np
2 import pickle
3
4 def softmax(X):
5     shifted_exp = np.exp(X.T - np.max(X, axis=1))
6     return (shifted_exp / np.sum(shifted_exp, axis=0)).T
7
8
9 def softmax_prime(output):
10    grad_start = np.einsum('ki, kj->kij', output, output)
11    identity = np.eye(grad_start.shape[1])
12    identities = np.stack([identity]*grad_start.shape[0])
13    output_matrix = np.stack([output]*output.shape[1], axis=2)
14    return np.multiply(output_matrix, identities) - grad_start
15
16
17 def sigmoid(X):
18     return 1/(1+np.exp(-X))
19
20
21 def sigmoid_prime(output):
22     return np.multiply(output, (1-output))
23
24
25 def cross_entropy_loss(X,y):
26     return -np.sum(np.multiply(y, np.log(X)))/y.shape[0]
27
28
29 def cross_entropy_prime(X,y):
30     return -y/(X)
31
32
33 def encode_words(indice_matrix, n_words=250):
34     indice_matrix = indice_matrix.T
35     if indice_matrix.shape[0] == 1:
36         return (np.arange(n_words) == indice_matrix[...,None]).astype(int)[0]
37     return (np.arange(n_words) == indice_matrix[...,None]).astype(int)
38
39
40 class Network:
41     def __init__(self, dictionary_size, embedding_size, hidden_layer_size, nwords=3,
42                  sigma=0.01, init_weights=True):
43         self.dict_size = dictionary_size
44         self.e_size = embedding_size
45         self.h_size = hidden_layer_size
46         self.n_words = nwords
```

```

47     if init_weights:
48         np.random.seed(3136)
49         self.W1 = np.random.normal(scale=sigma,size=(dictionary_size, embedding_size)
50                                     )
51
52         self.W2 = np.random.normal(scale=sigma, size=(embedding_size*nwords,
53                                     hidden_layer_size))
54         self.bias1 = np.random.normal(scale=sigma, size=(1,hidden_layer_size))
55
56         self.W3 = np.random.normal(scale=sigma,size=(hidden_layer_size,
57                                     dictionary_size))
58         self.bias2 = np.random.normal(scale=sigma, size=(1,dictionary_size))
59
60     def forward(self,words):
61
62         self.word1, self.word2, self.word3 = words
63
64         self.e1 = np.matmul(self.word1, self.W1)
65         self.e2 = np.matmul(self.word2, self.W1)
66         self.e3 = np.matmul(self.word3, self.W1)
67
68         self.e = np.concatenate([self.e1, self.e2, self.e3], axis=1)
69
70         self.h = np.matmul(self.e, self.W2) + self.bias1
71         self.h = sigmoid(self.h)
72
73         self.output = np.matmul(self.h, self.W3) + self.bias2
74         self.output = softmax(self.output)
75
76         return self.output
77
78     def backward(self, y):
79         X = self.output
80         batch_size = X.shape[0]
81
82         #self.loss_grad = cross_entropy_prime(X, y) #bs x 250
83         #self.softmax_grad = softmax_prime(X) # 250 x 250
84         #self.delta_3 = np.einsum("ki,kij->kj", self.loss_grad, self.softmax_grad) # bs
85         x 250
86
87         # To do things faster we can benefit from CE softmax derivative relation
88         delta_3 = X-y
89
90         self.W3_grad = np.matmul(delta_3.T, self.h).T / batch_size
91         self.bias2_grad = np.mean(delta_3, axis=0)
92
93         sigmoid_grad = sigmoid_prime(self.h)
94         delta_2 = np.multiply(np.dot(delta_3, self.W3.T), sigmoid_grad)
95
96         self.W2_grad = np.matmul(delta_2.T, self.e).T / batch_size
97         self.bias1_grad = np.mean(delta_2, axis=0)
98
99         delta_1 = np.matmul(delta_2, self.W2.T)

```

```

97
98     delta_1_1 = delta_1[:,self.e_size]
99     delta_1_2 = delta_1[:,self.e_size:self.e_size*2 ]
100    delta_1_3 = delta_1[:,self.e_size*2: ]
101
102
103    W1_grad1 = np.matmul(delta_1_1.T, self.word1).T / batch_size
104
105    W1_grad2 = np.matmul(delta_1_2.T, self.word2).T / batch_size
106
107    W1_grad3 = np.matmul(delta_1_3.T, self.word3).T / batch_size
108
109    self.W1_grad = W1_grad1 + W1_grad2 + W1_grad3
110
111    def save_network(self, save_path):
112
113        parameters = {}
114        parameters["W1"] = self.W1
115
116        parameters["W2"] = self.W2
117
118        parameters["W3"] = self.W3
119
120        parameters["bias2"] = self.bias2
121        parameters["bias1"] = self.bias1
122
123        with open(save_path, "wb") as file:
124            pickle.dump(parameters, file)
125
126    @classmethod
127    def load_network(cls, model_path):
128
129        with open(model_path, "rb") as file:
130            parameters = pickle.load(file)
131
132        cls.W1 = parameters["W1"]
133
134        cls.W2 = parameters["W2"]
135
136        cls.W3 = parameters["W3"]
137
138        cls.bias2 = parameters["bias2"]
139        cls.bias1 = parameters["bias1"]
140
141        dict_size = parameters["W1"].shape[0]
142        embedding_size = parameters["W1"].shape[1]
143        hidden_size = parameters["W2"].shape[1]
144        n = int(parameters["W2"].shape[0]/embedding_size)
145
146        return cls(dict_size, embedding_size, hidden_size, nwords=n, init_weights=False)
147
148
149    def SGD_step(self, lr=0.01):
150        self.W1 -= lr*self.W1_grad

```

```

151         self.W2 -= lr*self.W2_grad
152         self.W3 -= lr*self.W3_grad
153
154         self.bias1 -= lr*self.bias1_grad
155         self.bias2 -= lr*self.bias2_grad

```

train.py

```

1  import network as nn
2  import matplotlib.pyplot as plt
3  import numpy as np
4  import pickle
5  from time import time
6  import os
7
8  BATCH_SIZES = (16, 128, 512)
9  MAX_EPOCHS = 250
10 LRs= (0.1, 0.01,1)
11
12 checkpoint_paths = {}
13 max_no_improvement = 8
14
15 X_train = np.load('./data/train_inputs.npy')
16 y_train = np.load('./data/train_targets.npy').reshape(-1,1)
17
18 X_val = np.load('./data/valid_inputs.npy')
19 y_val = np.load('./data/valid_targets.npy').reshape(-1,1)
20
21 indices = list(range(X_train.shape[0]))
22 np.random.seed(3136)
23 np.random.shuffle(indices)
24 X_train = X_train[indices]
25 y_train = y_train[indices]
26
27 best_models = {}
28 for batch_size in BATCH_SIZES:
29     for lr in LRs:
30         save_name = f"bs_{batch_size}_lr_{lr}"
31         print(f"Training for parameters: Batch Size:{batch_size}, lr: {lr}")
32
33         n_batches = X_train.shape[0]//batch_size + ((X_train.shape[0]%batch_size !=0))
34
35         if not checkpoint_paths.get(save_name) is None:
36             network= nn.Network.load_network(checkpoint_paths.get(save_name))
37             print("Checkpoint Loaded...")
38         else:
39             network = nn.Network(250,16,128)
40
41         X_valid = nn.encode_words(X_val)
42         y_valid = nn.encode_words(y_val)
43
44         early_stop = 0
45         best_acc = 0
46         train_losses = []
47         validation_losses = []

```



```

48     train_accuracy = []
49     validation_accuracy = []
50     for epoch in range(MAX_EPOCHS):
51         start_time = time()
52         total_loss = 0
53         correct_words = 0
54         for i in range(n_batches):
55             start = i*batch_size
56             end = (i+1)*batch_size
57
58             X_batch = X_train[start:end]
59             X_batch = nn.encode_words(X_batch)
60
61             y_batch = y_train[start:end]
62             y_batch = nn.encode_words(y_batch)
63
64             y_pred = network.forward(X_batch)
65             network.backward(y_batch)
66
67             network.SGD_step(lr)
68
69             batch_loss = nn.cross_entropy_loss(y_pred, y_batch)
70             total_loss += batch_loss
71             correct_words += sum(np.argmax(y_pred, axis=1) == np.argmax(y_batch, axis
72                                     =1))
73
74             if (i % 1000) == 0:
75                 print(f"Epoch: {epoch}, Batch: {i} Train loss: {batch_loss}")
76
77         train_acc = correct_words/X_train.shape[0]
78         train_accuracy.append(train_acc)
79         train_losses.append(total_loss/n_batches)
80
81         print(f"Time: {time()-start_time} Validating Epoch {epoch}...")
82         y_pred = network.forward(X_valid)
83         loss = nn.cross_entropy_loss(y_pred, y_valid)
84         validation_losses.append(loss)
85         correct = sum(np.argmax(y_pred, axis=1) == np.argmax(y_valid, axis=1))
86         test_acc = correct/y_valid.shape[0]
87         validation_accuracy.append(test_acc)
88         print(f"Epoch {epoch}: Train loss: {train_losses[-1]}, Train accuracy:{
89                 train_acc}, Valid loss: {loss}, Valid accuracy: {test_acc}")
90         if not os.path.exists(f'./models/{save_name}/'):
91             os.makedirs(f'./models/{save_name}')
92         network.save_network(f'./models/{save_name}/Epoch{epoch}.pkl')
93         early_stop +=1
94         if test_acc >= best_acc:
95             early_stop = 0
96             best_acc = test_acc
97             network.save_network(f'./models/{save_name}/BestModel.pkl')
98         if early_stop > max_no_improvement:
99             break

```

```

100     metrics = {}
101     metrics["train_accuracy"] = train_accuracy
102     metrics["validation_accuracy"] = validation_accuracy
103     metrics["train_loss"] = train_losses
104     metrics["validation_loss"] = validation_losses
105     metrics["batch_size"] = batch_size
106     metrics["learning_rate"] = lr
107
108     with open(f"./models/summary_{save_name}.pkl", "wb") as file:
109         pickle.dump(metrics, file)
110
111     plt.style.use("ggplot")
112     plt.figure(figsize=(12,8))
113     plt.plot(list(range(len(train_accuracy))), train_accuracy, label='Train Accuracy')
114     plt.plot(list(range(len(validation_accuracy))), validation_accuracy, label='Validation Accuracy')
115     plt.xlabel("# of Epoch")
116     plt.ylabel("Accuracy")
117     plt.legend()
118     plt.savefig(f"./models/{save_name}/accuracy_plot.png")
119
120     plt.figure(figsize=(12,8))
121     plt.plot(list(range(len(train_losses))), train_losses, label='Train Loss')
122     plt.plot(list(range(len(validation_losses))), validation_losses, label='Validation Loss')
123     plt.xlabel("# of Epoch")
124     plt.ylabel("CE Loss")
125     plt.legend()
126     plt.savefig(f"./models/{save_name}/loss_plot.png")
127
128     best_models[save_name] = best_acc
129
130 with open("./models/best_models.pkl", "wb") as file:
131     pickle.dump(best_models, file)
132 print(best_models)
133 print("DONE")

```

eval.py

```

1 import network as nn
2 import numpy as np
3
4 MODEL_PATH = "./model.pkl"
5
6 X_train = np.load('./data/train_inputs.npy')
7 X_train = nn.encode_words(X_train)
8 y_train = np.load('./data/train_targets.npy').reshape(-1,1)
9 y_train = nn.encode_words(y_train)
10
11 X_val = np.load('./data/valid_inputs.npy')
12 X_val = nn.encode_words(X_val)
13 y_val = np.load('./data/valid_targets.npy').reshape(-1,1)
14 y_val = nn.encode_words(y_val)
15

```

```

16 X_test = np.load('./data/test_inputs.npy')
17 X_test = nn.encode_words(X_test)
18 y_test = np.load('./data/test_targets.npy').reshape(-1,1)
19 y_test = nn.encode_words(y_test)
20
21 vocab = list(np.load('./data/vocab.npy'))
22
23 network = nn.Network.load_network(MODEL_PATH)
24
25 city_of_new = 'city of new'
26 city_of_new_idx = [vocab.index(x) for x in city_of_new.split(' ')]
27
28 life_in_the = 'life in the'
29 life_in_the_idx = [vocab.index(x) for x in life_in_the.split(' ')]
30
31 he_is_the = 'he is the'
32 he_is_the_idx = [vocab.index(x) for x in he_is_the.split(' ')]
33
34 train_pred = network.forward(X_train)
35 valid_pred = network.forward(X_val)
36 test_pred = network.forward(X_test)
37
38 train_loss = nn.cross_entropy_loss(train_pred, y_train)
39 valid_loss = nn.cross_entropy_loss(valid_pred, y_val)
40 test_loss = nn.cross_entropy_loss(test_pred, y_test)
41
42 train_acc = sum(np.argmax(train_pred, axis=1) == np.argmax(y_train, axis=1))/y_train.
    shape[0]
43 valid_acc = sum(np.argmax(valid_pred, axis=1) == np.argmax(y_val, axis=1))/y_val.shape
    [0]
44 test_acc = sum(np.argmax(test_pred, axis=1) == np.argmax(y_test, axis=1))/y_test.shape
    [0]
45
46 batch = np.array([city_of_new_idx, life_in_the_idx, he_is_the_idx])
47 encoded_batch = nn.encode_words(batch)
48 pred_batch = network.forward(encoded_batch)
49 indices = np.argmax(pred_batch, axis=1)
50 words = [vocab[i] for i in indices]
51
52 print(f'Train Loss: {train_loss},    Train Accuracy: {train_acc}')
53 print(f'Validation Loss: {valid_loss}, Validation Accuracy: {valid_acc}')
54 print(f'Test Loss: {test_loss},      Test Accuracy: {test_acc}')
55
56 print("Predictions to examples:")
57 print(f"    {city_of_new} {words[0]}")
58 print(f"    {life_in_the} {words[1]}")
59 print(f"    {he_is_the} {words[2]}")

```

tsne.py

```

1 import network as nn
2 import numpy as np
3 from sklearn.manifold import TSNE
4 import matplotlib.pyplot as plt
5

```

```

6 MODEL_PATH = "./model.pkl"
7
8 vocab = list(np.load('./data/vocab.npy'))
9 network = nn.Network.load_network(MODEL_PATH)
10
11 encoded_words = np.identity(len(vocab))
12 embeddings = np.matmul(encoded_words, network.W1)
13
14 tsne_model = TSNE(perplexity=40, n_components=2, init='pca', n_iter=2500, random_state
    =3136)
15 reduced_embeddings = tsne_model.fit_transform(embeddings)
16
17 x = reduced_embeddings[:,0]
18 y = reduced_embeddings[:,1]
19
20 plt.style.use("ggplot")
21 plt.figure(figsize=(12,12))
22 plt.scatter(x, y)
23 for label, xi, yi in zip(vocab, x, y):
24     plt.annotate(label, xy=(xi,yi), xytext=(0, 0), textcoords='offset points')
25
26 plt.xlabel("Dim 1")
27 plt.ylabel("Dim 2")
28 plt.title("TSNE Plot of Reduced Embeddings")
29
30 plt.savefig('./tsne_plot.png')
31 print("Plot saved to directory")

```